

# UDACITY MACHINE LEARNING PROJECT CAPSTONE

June 2021

## 1. Project overview

The project goal is to build a Dog Breeds Classifier app that uses an image classification deep learning model to perform the Dog breed identification. To achieve this goal is necessary to review different models based on CNN architectures (Convolutional Neuronal Networks) mainly. The project development includes to try models built from scratch and models using transfer learning approach on pretrained image classification architectures. The main tool used to develop the project is Pytorch and the models available in the torchvision.models library.

## 2. Project description

Project starts with a data exploration to review the image data organization and classes defined. Data formats need to be transformed to fit the pretrained model's requirements building data loaders function that process the images and add data augmentation to improve the training process. Different libraries and frameworks are suggested to use for the task required in the app. Face detection, Dog's detection, and Dog Breed's detection. The app to be developed must perform these tasks:

- If a dog is detected in the image, return the predicted breed.
- If a human is detected in the image, return the resembling dog breed.
- If neither is detected in the image, provide output that indicates an error.


Some models are suggested by the Udacity Team, but other models will be considered and tested to achieve best app performance.

## 3. Data exploration

Dataset used for the training process consist in face images and dog's images classified by name and dog breed, respectively. There are 13.233 human images and 8.351 dog images. For the Dogs breed classifier, the dog images provided by Udacity are organized by train, test and validation sets as subfolders.

001.Affenpinscher	056.Dachshund	111.Norwich_terrier	025.Black_and_tan_coonhound	080.Greater_swiss_mountain_dog
002.Afghan_hound	057.Dalmatian	112.Nova_scotia_duck_tolling_retriever	026.Black_russian_terrier	081.Greyhound
003.Airedale_terrier	058.Dandie_dinmont_terrier	113.Old_english_sheepdog	027.Bloodhound	082.Havanese
004.Akita	059.Doberman_pinscher	114.Otterhound	028.Blue_tick_coonhound	083.Ibizan_hound
005.Alaskan_malamute	060.Dogue_de_bordeaux	115.Papillon	029.Border_collie	084.Icelandic_sheepdog
006.American_eskimo_dog	061.English_cocker_spaniel	116.Parson_russell_terrier	030.Border_terrier	085.Irish_red_and_white_setter
007.American_foxhound	062.English_setter	117.Pekingese	031.Borzoi	086.Irish_setter
008.American_staffordshire_terrier	063.English_springer_spaniel	118.Pembroke_welsh_corgi	032.Boston_terrier	087.Irish_terrier
009.American_water_spaniel	064.English_toy_spaniel	119.Petit_basset_griffon_vendéen	033.Bouvier_des_flandres	088.Irish_water_spaniel
010.Anatolian_shepherd_dog	065.Entlebucher_mountain_dog	120.Pharao_hound	034.Boxer	089.Irish_wolfhound
011.Australian_cattle_dog	066.Field_spaniel	121.Plott	035.Boykin_spaniel	090.Italian_greyhound
012.Australian_shepherd	067.Finnish_spitz	122.Pointer	036.Briard	091.Japanese_chin
013.Australian_terrier	068.Flat-coated_retriever	123.Pomeranian	037.Brittany	092.Keeshond
014.Basenji	069.French_bulldog	124.Poodle	038.Brussels_griffon	093.Kerry_blue_terrier
015.Basset_hound	070.German_pinscher	125.Portuguese_water_dog	039.Bull_terrier	094.Komondor
016.Beagle	071.German_shepherd_dog	126.Saint_bernard	040.Bulldog	095.Kuvasz
017.Bearded_collie	072.German_shorthaired_pointer	127.Silky_terrier	041.Bullmastiff	096.Labrador_retriever
018.Beauceron	073.German_wirehaired_pointer	128.Smooth_fox_terrier	042.Cairn_terrier	097.Lakeland_terrier
019.Bedlington_terrier	074.Giant_schnauzer	129.Tibetan_mastiff	043.Canaan_dog	098.Leonberger
020.Belgian_malinois	075.Glen_of_imaal_terrier	130.Welsh_springer_spaniel	044.Cane_corso	099.Lhasa_apso
021.Belgian_sheepdog	076.Golden_retriever	131.Wirehaired_pointing_griffon	045.Cardigan_welsh_corgi	100.Lowchen
022.Belgian_tervuren	077.Gordon_setter	132.Xoloitzcuintli	046.Cavalier_king_charles_spaniel	101.Malte
023.Bernese_mountain_dog	078.Great_dane	133.Yorkshire_terrier	047.Chesapeake_bay_retriever	102.Manchester_terrier
024.Bichon_frise	079.Great_pyrenees		048.Chihuahua	103.Mastiff
			049.Chinese_crested	104.Minature_schnauzer
			050.Chinese_shar-pei	105.Neapolitan_mastiff
			051.Chow_chow	106.Newfoundland
			052.Clumber_spaniel	107.Norfolk_terrier
			053.Cocker_spaniel	108.Norwegian_buhund
			054.Collie	109.Norwegian_elkhound
			055.Curly-coated_retriever	110.Norwegian_Lundehund

There are 133 breeds categories where Border Collie is the most frequent class and Xoloitscuintli the less one.



Images are in RGB format and different sizes. These are an images sample:



## 4. Metrics

**i** Metric used for model training was accuracy to assess the model classification quality. The accuracy formula is:

$$Accuracy = \frac{TrueNegatives + TruePositive}{TruePositive + FalsePositive + TrueNegative + FalseNegative}$$

This is a percentage between 0 to 1. Higher values mean better classification quality of the model.

## 5. Data Pre-processing

The input images have different sizes, for that reason is necessary do some pre-processing. Using torchvision.transforms from PyTorch framework these transformation were applied to the data input:

- `transforms.Resize(IMAGE_RESIZE)`: Resized the images to 256 px as is required by the Image classification Models.
- `transforms.RandomRotation(28)`: Apply a random rotation between 0-28 grades
- `transforms.RandomHorizontalFlip()`: Generate random horizontal flip for data augmentation
- `transforms.CenterCrop((IMAGE_SIZE, IMAGE_SIZE))`: Center cropped to 224 px
- `transforms.Normalize`: Apply normalization with 0.5 mean and 0.5 std

## 6. Reference architectures

Reviewing the state-of-the-art image classification models, the site presents the recent models proposed and the accuracy achieved using the ImageNet dataset.


<https://paperswithcode.com/sota/image-classification-on-imagenet>. The models available in the standard `torchvision.models` (Last Torchvision version 0.8):

- AlexNet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3
- GoogLeNet
- ShuffleNet v2
- MobileNetV2
- MobileNetV3
- ResNeXt
- Wide ResNet
- MNASNet

Models available that perform better are:

- **resnext101\_32x8d**: The model is the same as ResNet except for the bottleneck number of channels which is twice larger in every block.
- **resnet152**: Residual Networks with 152 layers.
- **Inception v3**: convolutional neural network architecture from the Inception family that makes several improvements including using Label Smoothing, factorized 7 x 7 convolutions, and the use of an auxiliary classifier to propagate label information lower down the network (along with the use of batch normalization for layers in the side head).

Resnext101 uses over 80 million parameters, Inception v3 require over 20 million parameters, but state-of-the-art that performs better on the Imagenet dataset uses over 600 million parameters, that can use more time, inclusive for retraining the last layer for transfer learning. For that reason, I am reviewing the models pre-trained in Torchvision and validate that best performing model is resnext101\_32x8d, but this one is only available in torchvision 0.8.



Models with to performance are:

Rank	Model	Top 1 Accuracy ↑	Top 5 Accuracy	Number of params	extra Training Data	Paper	Code	Result	Year
1	Meta Pseudo Labels (EfficientNet-L2)	90.2%	98.8%	480M	✓	Meta Pseudo Labels	<a href="#">GitHub</a>	<a href="#">Result</a>	2021
2	Meta Pseudo Labels (EfficientNet-B6-Wide)	90%	98.7%	390M	✓	Meta Pseudo Labels	<a href="#">GitHub</a>	<a href="#">Result</a>	2021
3	NFNet-F4+	89.2%		527M	✓	High-Performance Large-Scale Image Recognition Without Normalization	<a href="#">GitHub</a>	<a href="#">Result</a>	2021
4	ALIGN (EfficientNet-L2)	88.64%	98.67%	480M	✓	Scaling Up Visual and Vision-Language Representation Learning With Noisy Text Supervision	<a href="#">GitHub</a>	<a href="#">Result</a>	2021
5	EfficientNet-L2-475 (SAM)	88.61%		480M	✓	Sharpness-Aware Minimization for Efficiently Improving Generalization	<a href="#">GitHub</a>	<a href="#">Result</a>	2020
6	VIT-H/14	88.55%		632M	✓	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale	<a href="#">GitHub</a>	<a href="#">Result</a>	2020
7	FixEfficientNet-L2	88.5%	98.7%	480M	✓	Fixing the train-test resolution discrepancy: FixEfficientNet	<a href="#">GitHub</a>	<a href="#">Result</a>	2020
8	NoisyStudent (EfficientNet-L2)	88.4%	98.7%	480M	✓	Self-training with Noisy Student improves ImageNet classification	<a href="#">GitHub</a>	<a href="#">Result</a>	2020
9	Mixer-H/14 (JFT-300M pre-train)	87.94%			✓	MLP-Mixer: An all-MLP Architecture for Vision	<a href="#">GitHub</a>	<a href="#">Result</a>	2021

resnext101\_32x8d looks like the best performance model in the list of available models, but still far from the lasted state-of-the-art models. Despite requires over 80 million parameters. I am trying to use this model for the transfer learning process. I decided to use resnet\_150 because the number of parameters still low and is better for the resources available. This model requires image input with size 224x224 and has up to 2040 output classes. I also added to layers more for the transfer learning.

## 7. Model experimentation from scratch

First, I read some references about CNN for image classification with some examples and recommendations. The main reference used was: <https://machinelearningmastery.com/how-to-develop-a-cnn-from-scratch-for-cifar-10-photo-classification/>. Some recommendation is to define different models with different complexity (Layer numbers and pooling operations) and check for normalization, regularization, and dropout to improve results and avoid overfitting. All models get an input of 224 x 224 x 3 (Image size and RGB channels) and output for 133 classes. I defined three different models:

### Base model (Small)

This is the base model, which has basic convolutional layers and a fully connected layer to map the tensors to the output size.

- 3 convolutional layers
- Max pooling
- Batch normalization
- Dropout
- 2 fully connected layers

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 112, 112]	896
MaxPool2d-2	[-1, 32, 56, 56]	0
Conv2d-3	[-1, 64, 28, 28]	18,496
MaxPool2d-4	[-1, 64, 14, 14]	0
Conv2d-5	[-1, 32, 14, 14]	18,464
MaxPool2d-6	[-1, 32, 7, 7]	0
Dropout-7	[-1, 1568]	0
Linear-8	[-1, 133]	208,677
BatchNorm1d-9	[-1, 133]	266

- Total params: 910,437
- Trainable params: 910,437
- Non-trainable params: 0

For this model was used 30 epochs

### Medium size model

This model a couple more convolutional layers trying to improve the classification task and add batch normalization between convolutional layers.

- 5 convolutional layers
- Max pooling
- Batch normalization
- Dropout
- 3 fully connected layers

Layer (type)	Output Shape	Param #
--------------	--------------	---------

Conv2d-1	[-1, 128, 112, 112]	3,584
BatchNorm2d-2	[-1, 128, 112, 112]	256
MaxPool2d-3	[-1, 128, 56, 56]	0
Conv2d-4	[-1, 64, 56, 56]	73,792
BatchNorm2d-5	[-1, 64, 56, 56]	128
Conv2d-6	[-1, 32, 56, 56]	18,464
MaxPool2d-7	[-1, 32, 28, 28]	0
Conv2d-8	[-1, 64, 28, 28]	18,496
MaxPool2d-9	[-1, 64, 14, 14]	0
Conv2d-10	[-1, 32, 14, 14]	18,464
MaxPool2d-11	[-1, 32, 7, 7]	0
Dropout-12	[-1, 1568]	0
Linear-13	[-1, 784]	1,230,096
Dropout-14	[-1, 784]	0
Linear-15	[-1, 382]	299,870
Dropout-16	[-1, 382]	0
Linear-17	[-1, 133]	50,939
BatchNorm1d-18	[-1, 133]	266

- Total params: 1,714,355
- Trainable params: 1,714,355
- Non-trainable params: 0

### Big size model

This mode adds more convolutional layer trying to capture more image characteristics but increase the number of the parameters over 3 million making the model take more than 10 hours to be trained. The model is basically like the base model but adds 7 more convolutional layers.

- 8 convolutional layers
- Max pooling
- Batch normalization
- Dropout
- Fully connected layers

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
BatchNorm2d-2	[-1, 64, 224, 224]	128
Conv2d-3	[-1, 128, 224, 224]	73,856
BatchNorm2d-4	[-1, 128, 224, 224]	256
Conv2d-5	[-1, 128, 224, 224]	147,584
MaxPool2d-6	[-1, 128, 112, 112]	0
Conv2d-7	[-1, 256, 112, 112]	295,168
BatchNorm2d-8	[-1, 256, 112, 112]	512
Conv2d-9	[-1, 256, 112, 112]	590,080
MaxPool2d-10	[-1, 256, 56, 56]	0
Conv2d-11	[-1, 128, 56, 56]	295,040
MaxPool2d-12	[-1, 128, 28, 28]	0
Conv2d-13	[-1, 64, 28, 28]	73,792
MaxPool2d-14	[-1, 64, 14, 14]	0
Conv2d-15	[-1, 32, 14, 14]	18,464
MaxPool2d-16	[-1, 32, 7, 7]	0
Linear-17	[-1, 784]	1,230,096
Dropout-18	[-1, 784]	0

Linear-19	[-1, 392]	307,720
Dropout-20	[-1, 392]	0
Linear-21	[-1, 133]	52,269
BatchNorm1d-22	[-1, 133]	266

- Total params: 3,087,023
- Trainable params: 3,087,023
- Non-trainable params: 0

For this model, the result running only 20 epochs executed due to the long training time required.

After try the different models, there results were obtained:

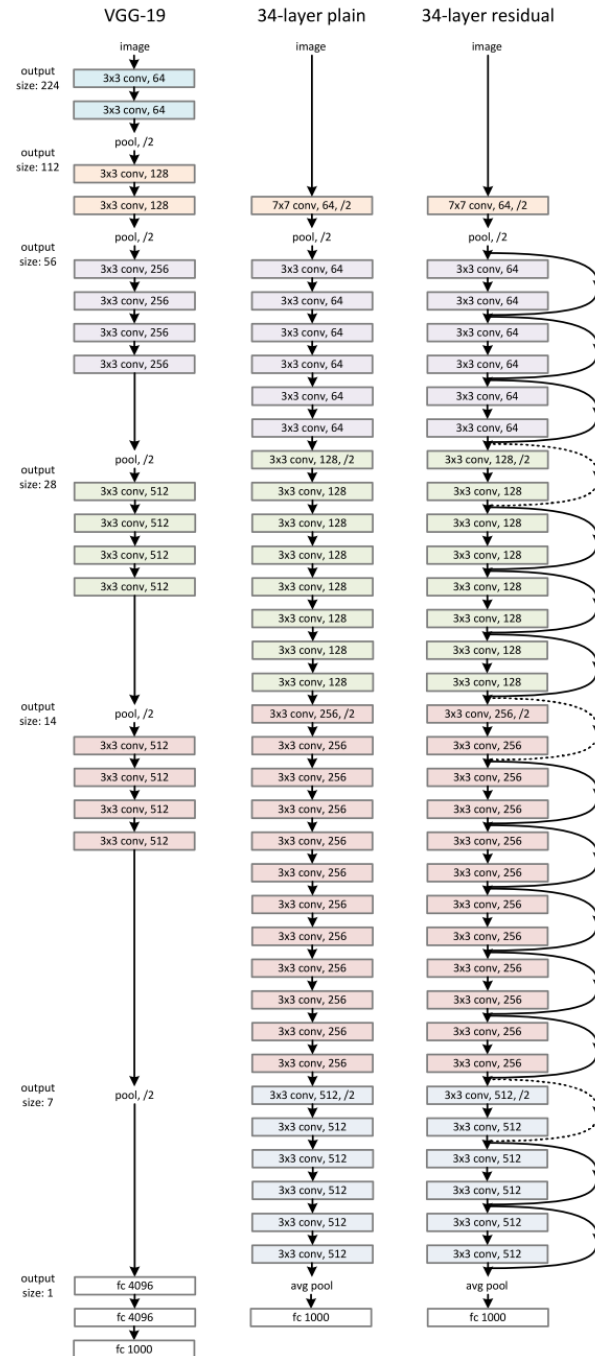
Model	Test Loss	Parameters number	Epochs	Test Accuracy
Small size	3.232469	246,799	50	21% (183/836)
Medium size	3.131659	1,714,355	50	25% (215/836)
Big size	3.917502	3,087,023	20	11% (94/836)

## 8. Model implementation using Transfer Learning.

To improve the results, transfer Learning approach was used using pretrained models provide by torchvision library. I tried different models like Resnet-50 and Inception v2 but the accuracy was not the best and moved to Resnet152. I needed to add an additional layer for the transfer learning and with different Learning rates and optimizers. Finally, the best performance model achieved over 85% accuracy with 15 epochs.

The resnet152 architecture has 59,261,125 training parameters and looks like this:





Two fully connected layers were added to the model to perform the final classification with the breed samples. Relu activation functions was used and dropout regularization with 0.3.

```

import torchvision.models as models
import torch.nn as nn
from collections import OrderedDict

## TODO: Specify model architecture
INPUT_FEATURES = 2048
OUTPUT_FEATURES = 133

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# Load the model from torch models to perform the transfer learning using inception_v3
model_transfer_learning = models.resnet152(pretrained=True)

# freeze all model parameters
for parameter in model_transfer_learning.parameters():
    parameter.requires_grad = False

# Get the input features size for the model resnet 152
input_features = model_transfer_learning.fc.in_features

# This is recommended to run deep models to stabilize the training process
# https://discuss.pytorch.org/t/why-auxiliary-Logits-set-to-false-in-train-mode/40705
model_transfer_learning.aux_logits=False

# Define the additional Layer for the transfer Learning process
model_transfer_learning.fc = nn.Sequential(OrderedDict(
    [
        ('fc1', nn.Linear(INPUT_FEATURES, 512)),
        ('relu', nn.ReLU()),
        ('dropout', nn.Dropout(p=0.3)),
        ('fc2', nn.Linear(512, OUTPUT_FEATURES))
    ])

# Verify if cuda GPU is available
if use_cuda:
    model_transfer_learning = model_transfer_learning.cuda()

```

## Loss Function and Optimizer

```

# criterion_transfer = None
# optimizer_transfer = None

# Define the Learning rate parameter
LEARNING_RATE = 0.02

# Define the Loss function
criterion_transfer_learning = nn.CrossEntropyLoss()

# Define the optimizer
# Adam optimizer
optimizer_transfer_learning = optim.SGD(model_transfer_learning.fc.parameters(), lr=LEARNING_RATE)

```

## 9. App implementation and results

With the final model trained using transfer learning, a prediction function was implemented to use in the final app.

### Prediction function:

```
def predict_breed_transfer(img_path):
    # Load the image and return the predicted breed

    # Open the image related to the img_path
    img=Image.open(img_path).convert("RGB")

    # Define image transformation
    img_transform = transforms.Compose([transforms.Resize(256),
                                       transforms.CenterCrop(224),
                                       transforms.ToTensor(),
                                       transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                             std=[0.229, 0.224, 0.225])
                                       ])

    # Create the reshaped image tensor
    image_tensor = img_transform(img).unsqueeze(0)

    # Validate the cuda availability
    if use_cuda:
        image_tensor = image_tensor.cuda()

    # Performing inference
    model_transfer_learning.eval()


    # Execute the inference and apply softmax function and extract the final value
    output = F.softmax(model_transfer_learning(image_tensor),dim=1).cpu().data.numpy().squeeze()

    return class_names[np.argmax(output)]
```

### Testing function:

```
# Test the function
dog_files_test = np.array(glob("/data/dog_images/test/*/*"))
image_index = int(np.random.randint(0,len(dog_files_test),1))
breed_detected = predict_breed_transfer(dog_files_test[image_index])
print("Image file Input : {}",dog_files_test[image_index])
image = Image.open(dog_files_test[image_index])
plt.imshow(image)
plt.show()
print('Predicted dog breed: {}'.format(breed_detected))
```

Image file Input : {} /data/dog\_images/test/046.Cavalier\_king\_charles\_spaniel/Cavalier\_king\_charles\_spaniel\_03258.jpg




Predicted dog breed: {} Cavalier king charles spaniel

## Final App implementation

```
def run_app(img_path):  
    '''  
    Define a function to execute inference process to detect faces and fogd bred  
    predicted ImageNet class for image at specified path  
  
    Args:  
        img_path: path to an image  
  
    Returns:  
        Index corresponding to RESNET-120 model's prediction  
    '''  
  
    ## handle cases for a human face, dog, and neither  
  
    # First Look for human face  
    if face_detector(img_path):  
        print("Face Detected")  
        image = Image.open(img_path)  
        plt.imshow(image)  
        plt.show()  
        print(f"Looks like: \n The {predict_breed_transfer(img_path)}\n")  
  
    # If not face detectec check for dog breed  
    elif dog_detector(img_path):  
        print("Dog Detected")  
        image = Image.open(img_path)  
        plt.imshow(image)  
        plt.show()  
        print(f"Predicted breed:... \n{predict_breed_transfer(img_path)}\n")  
  
    # Not face or dog detected  
    else:  
        print("Couldn't detect dogs or faces.\n")  
        image = Image.open(img_path)  
        plt.imshow(image)  
        plt.show()
```

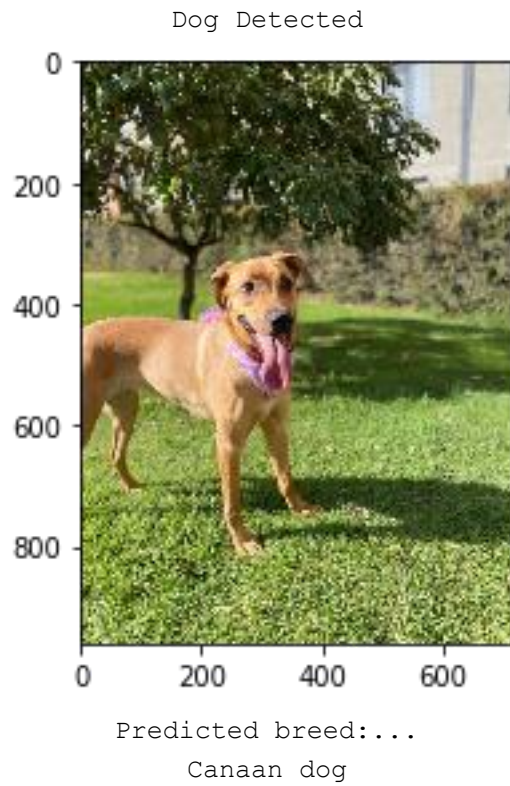
## Results

### Example 1

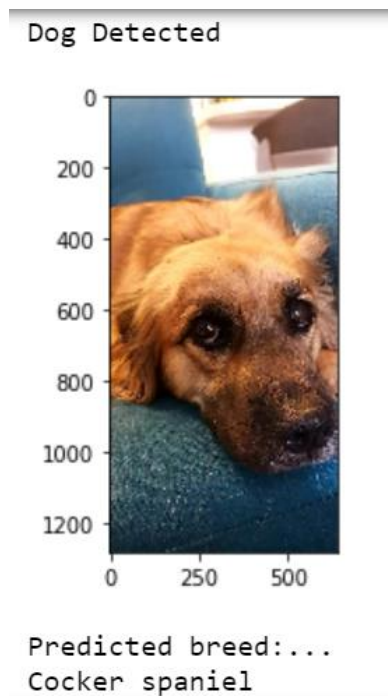


In this case the classifier failed, detecting a human face from the dog picture.

### Example 2



### Example 3




#### Example 4

Couldn't detect dogs or faces.




Bad luck for the Grinch with the breed dog classifier


#### Example 5



#### Example 6



Looks like:  
The Black russian terrier  
Smart classifier!



My Shih tzu doggy coco is sad for my classifier results!

## 10. Conclusion

The model trained using transfer learning with resnet152 pre-trained model achieved around 85% accuracy. I tried with some additional images, for example, two ShihZu images, and the model does not work well. and with dogs with mixed bred the model was confused. That means that there is a lot of improvement because 85% accuracy sounds ok but in practice still far from a good performance. Is clear that using a pre-trained model offers a fast and better way to train models for a specific task like dog breed detection. As I suggested when chose the pre-trained model, there more state-of-the-art models with more than 600 million parameters.

### *Possible improvements*

1. As was reviewed before, state-of-the-art models offer higher accuracy for the image classification tasks, can also provide better performance for the transfer learning task. These models require more resources for the training process.
2. Additional optimizations can be used to improve performance, like hyperparameter tuning and network pruning to accelerate inference time response.

3. Increasing data with more sources using web scrapping and data augmentation using GAN networks to enrich the training process can provide better input for the training and validation process.
4. New models based on a generative approach can be used to improve the transfer learning process.